

Confidential Containers

Chris Butler
Chief Architect, CTO Organization
chris.butler@redhat.com


Dave Ripani
Solution Architect
rippa@redhat.com

Agenda


- ▶ Confidential Computing 101
- ▶ Confidential containers architecture
- ▶ Demo 1: Protecting data in memory
- ▶ RATS architecture
- ▶ Use case scenarios

Focus areas for confidential computing


Multi-level security



Compartmented data



Physical asset
compromise



Demo 1: - Showing it protects memory

- Scenario 1: Standard containers (I go to root in the box and find it)
- Scenario 2: kata containers - (e.g. shim qemu) - get in as root; use gcore / find string
- Scenario 3: 'CoCo' container - go in dump memory show we can't find it.



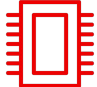
Demo 2:

- **Attestation and releasing secret**
 - Showing in a TDX container I can get the secret otherwise I can't

Deployment arch (high level)

- Multicluster
-

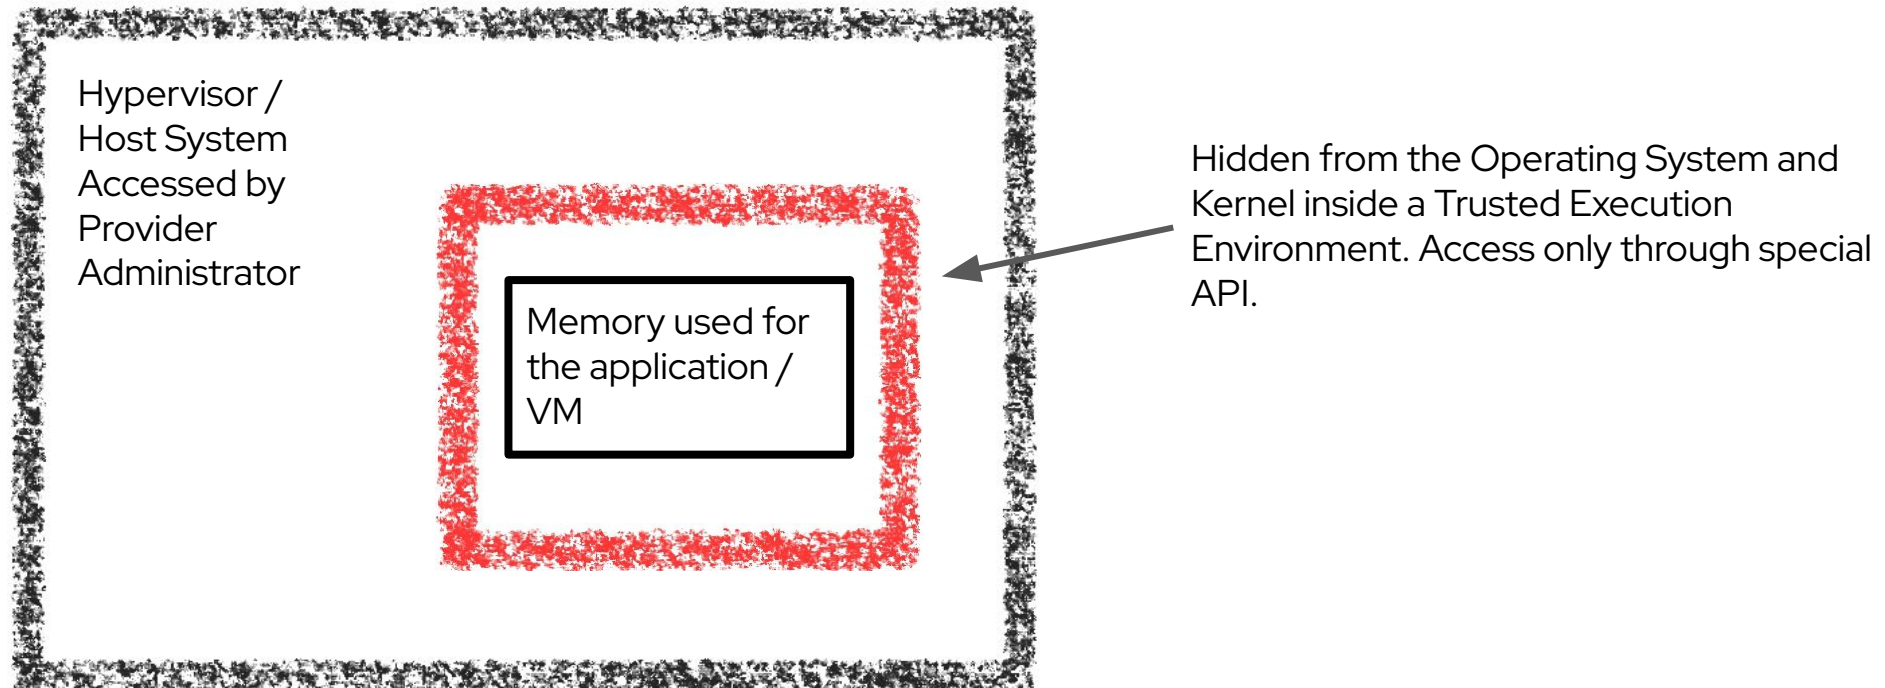
Confidential computing enforces cryptographic protection of applications and data

○ storing		→ at-rest	LUKS, GnuPG,...
○ transmitting		→ in-transit	OpenSSL, TLS,...
○ processing		→ in-use	<div>Confidential Computing is about protecting data in-use. (incl. integrity)</div>

Enforcing that you do not trust the system admins or infra provider admins

Data in use

- ▶ How does Confidential Computing work?

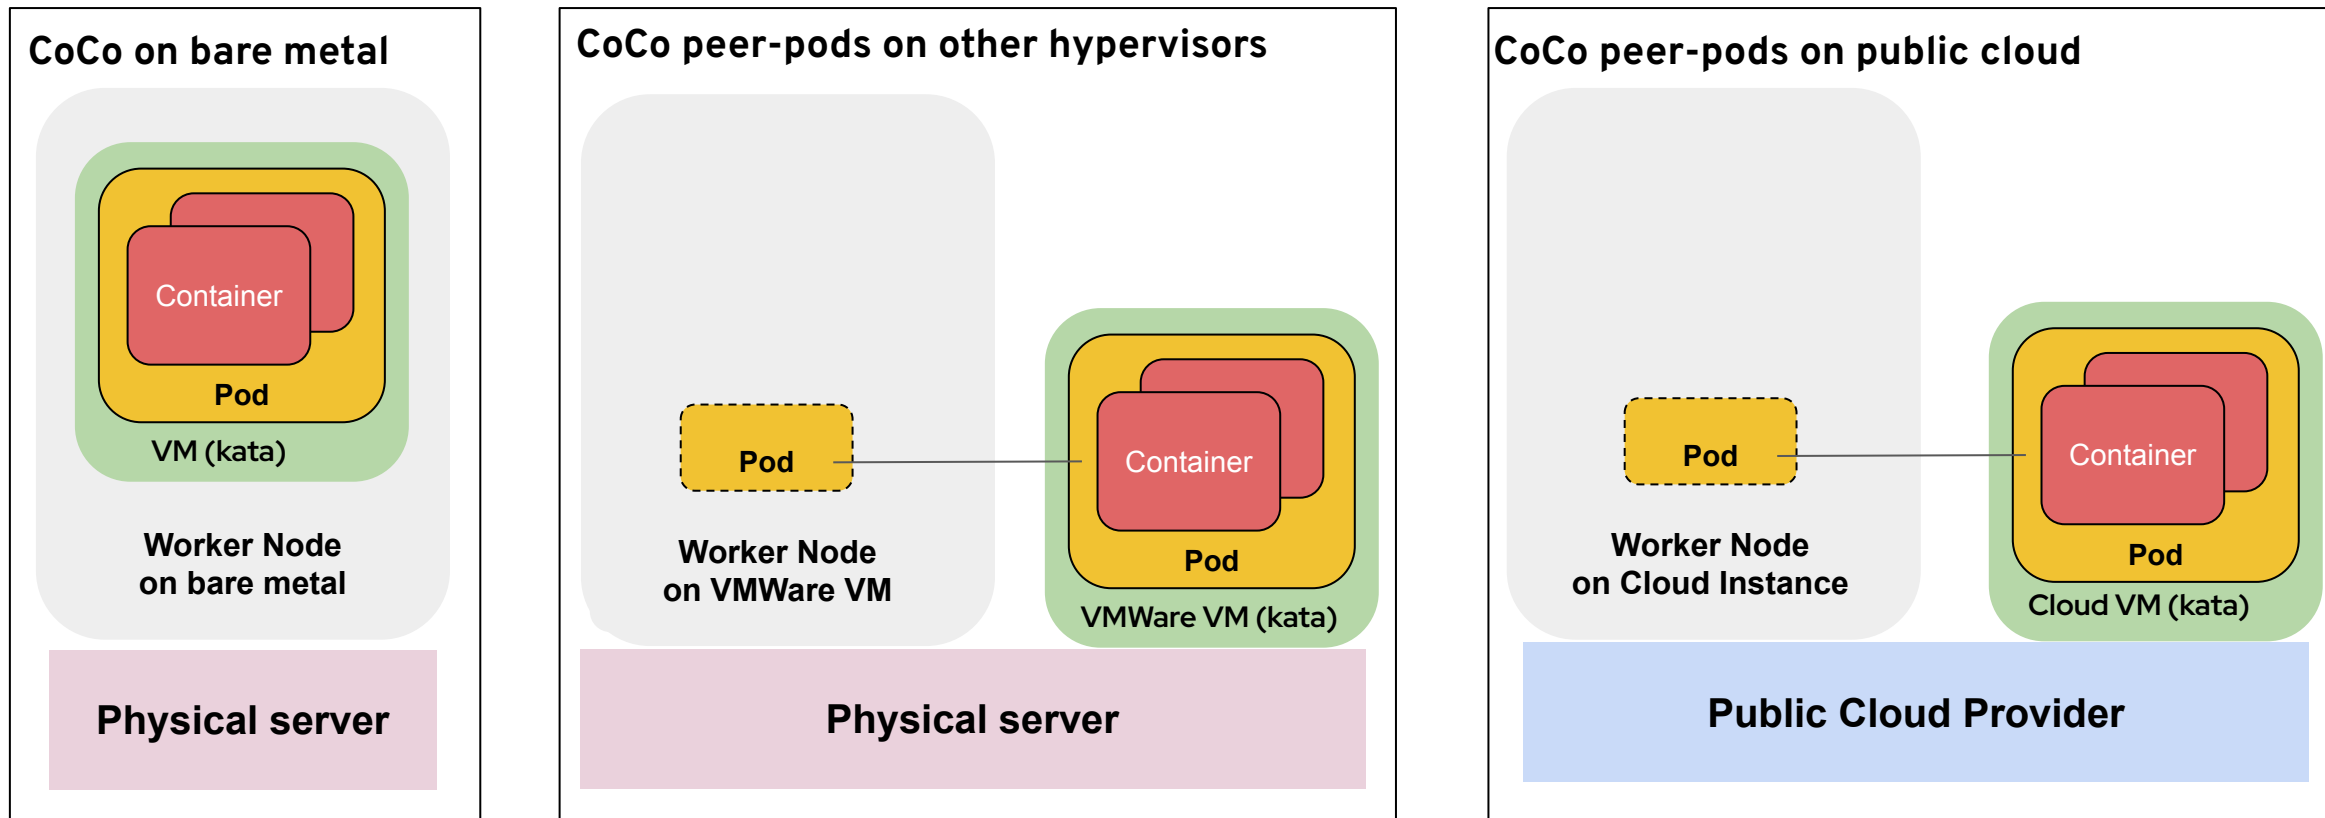


Confidential Containers (CNCF)



Confidential Containers is an **open source community** working to **enable cloud native confidential computing** by leveraging Trusted Execution Environments to protect containers and data.

Confidential Containers across cloud & bare metal



- ▶ Peer-pods evolves the OSC solution from bare metal to hypervisors and public cloud
- ▶ Peer-pods code repositories enhance the [kata containers](#) project and the [confidential containers](#) project

Demo one

Memory isolation in confidential containers

Select a CoCo runtime

Pull data a 'secret' in

Do nothing

No special privileges

Store the secret in memory

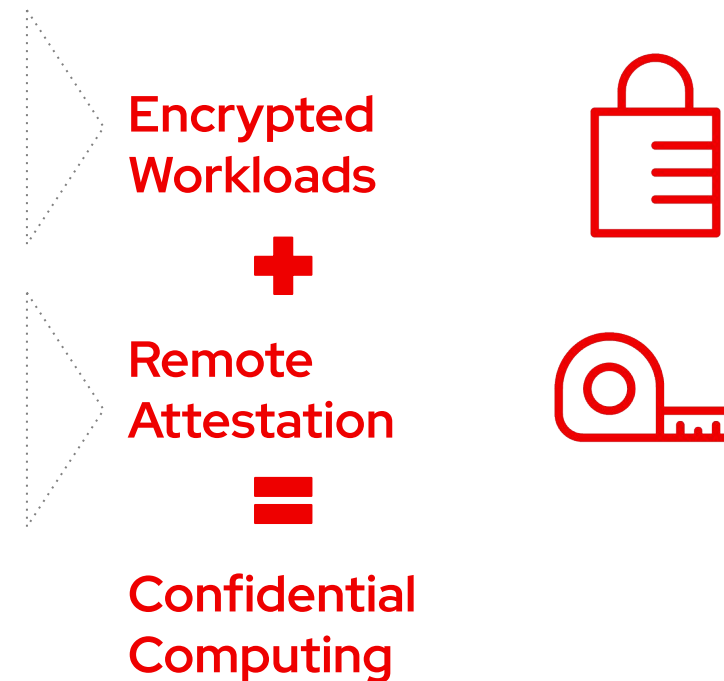
```
apiVersion: v1
kind: Pod
metadata:
  name: coco-demo
  annotations:
    io.katacontainers.config.hypervisor.default_vcpus: "2"
    io.katacontainers.config.hypervisor.default_memory: "4096"
spec:
  runtimeClassName: kata-cc-tdx
  #runtimeClassName: kata
  initContainers:
    - name: fetch-key
      image: registry.access.redhat.com/ubi9/ubi:9.3
      command:
        - sh
        - -c
        - curl -L https://gist.githubusercontent.com/butler54/21a0.../raw.txt -o /keys/magickey.txt
      volumeMounts:
        - name: keys
          mountPath: /keys
  containers:
    - name: coco-demo
      image: registry.access.redhat.com/ubi9/ubi:9.3
      command:
        - sleep
        - "36000"
      securityContext:
        privileged: false
        seccompProfile:
          type: RuntimeDefault
      volumeMounts:
        - name: keys
          mountPath: /keys
  volumes:
    - name: keys
      emptyDir:
        medium: Memory
```

Threat Vectors:

Pod Images: Risk of tampering/access.

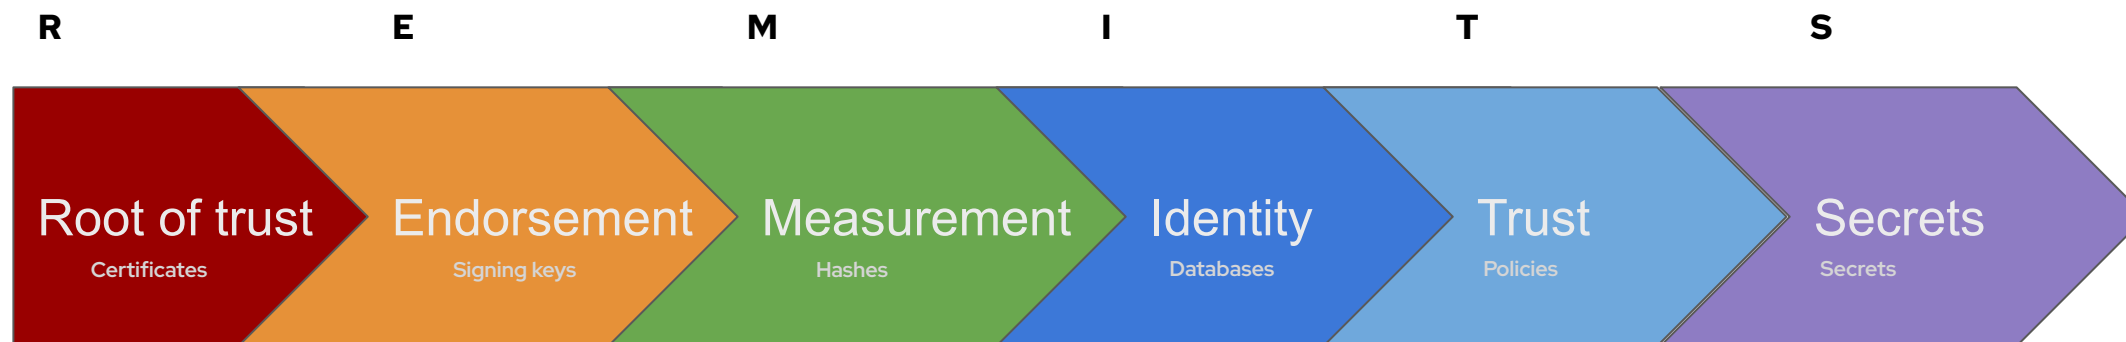
Pod Memory: Infrastructure provider access.

Pod Data: Provider tampering/access.



Establishing a chain of trust

REMITs



Example:

private key from a
chip manufacturer



per device
authenticity info

crypto hash over
code & config

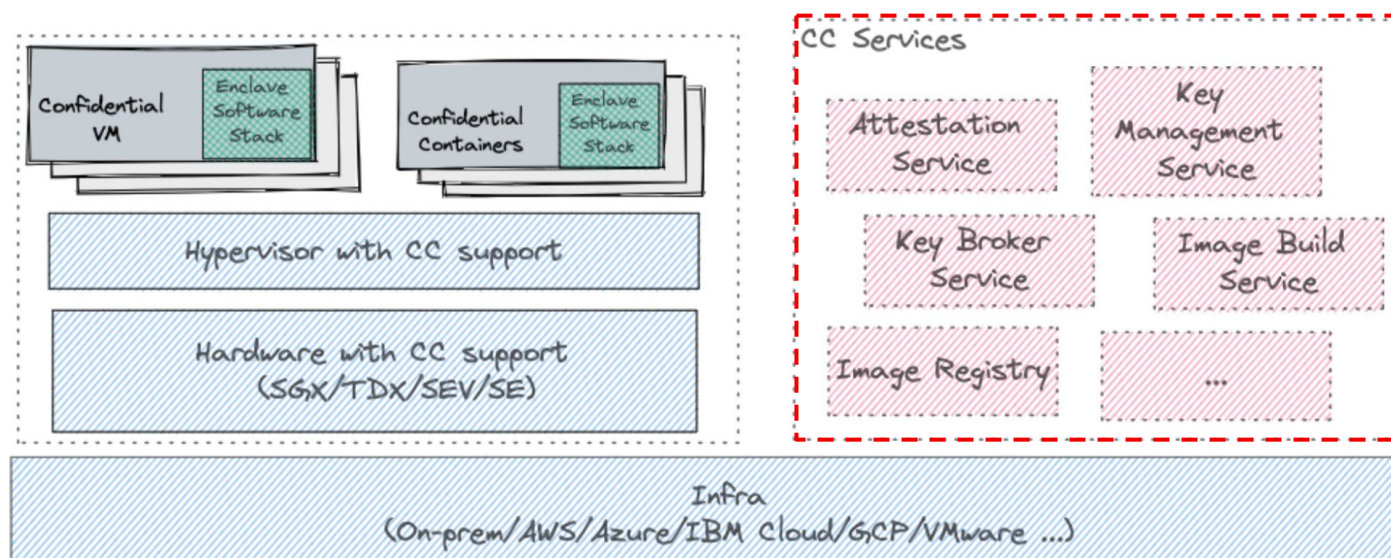
compare evidence
with ref. values

apply diff. policies
for diff. envs

share secrets, e.g.
encryption keys



Confidential Computing Services




- **Attestation Service** – The primary purpose of the attestation service is to validate the evidence provided by the hardware TEE. This is the *Verifier*, as defined in the [RATS architecture](#).
- **Key Broker Service (KBS)** – The KBS is the *Relying Party*, as defined by the [RATS architecture](#). Following are its primary functions:
 - Receive evidence from the *Attester* (confidential VM or container) via a challenge-response protocol.
 - Relay the evidence to the Attestation Service for verification.
 - Apply appraisal policy for the returned Attestation Results to assess the trustworthiness of the *Attester*.
 - Interact with the Key Management Service to retrieve the keys and then send them back to the *Attester*.

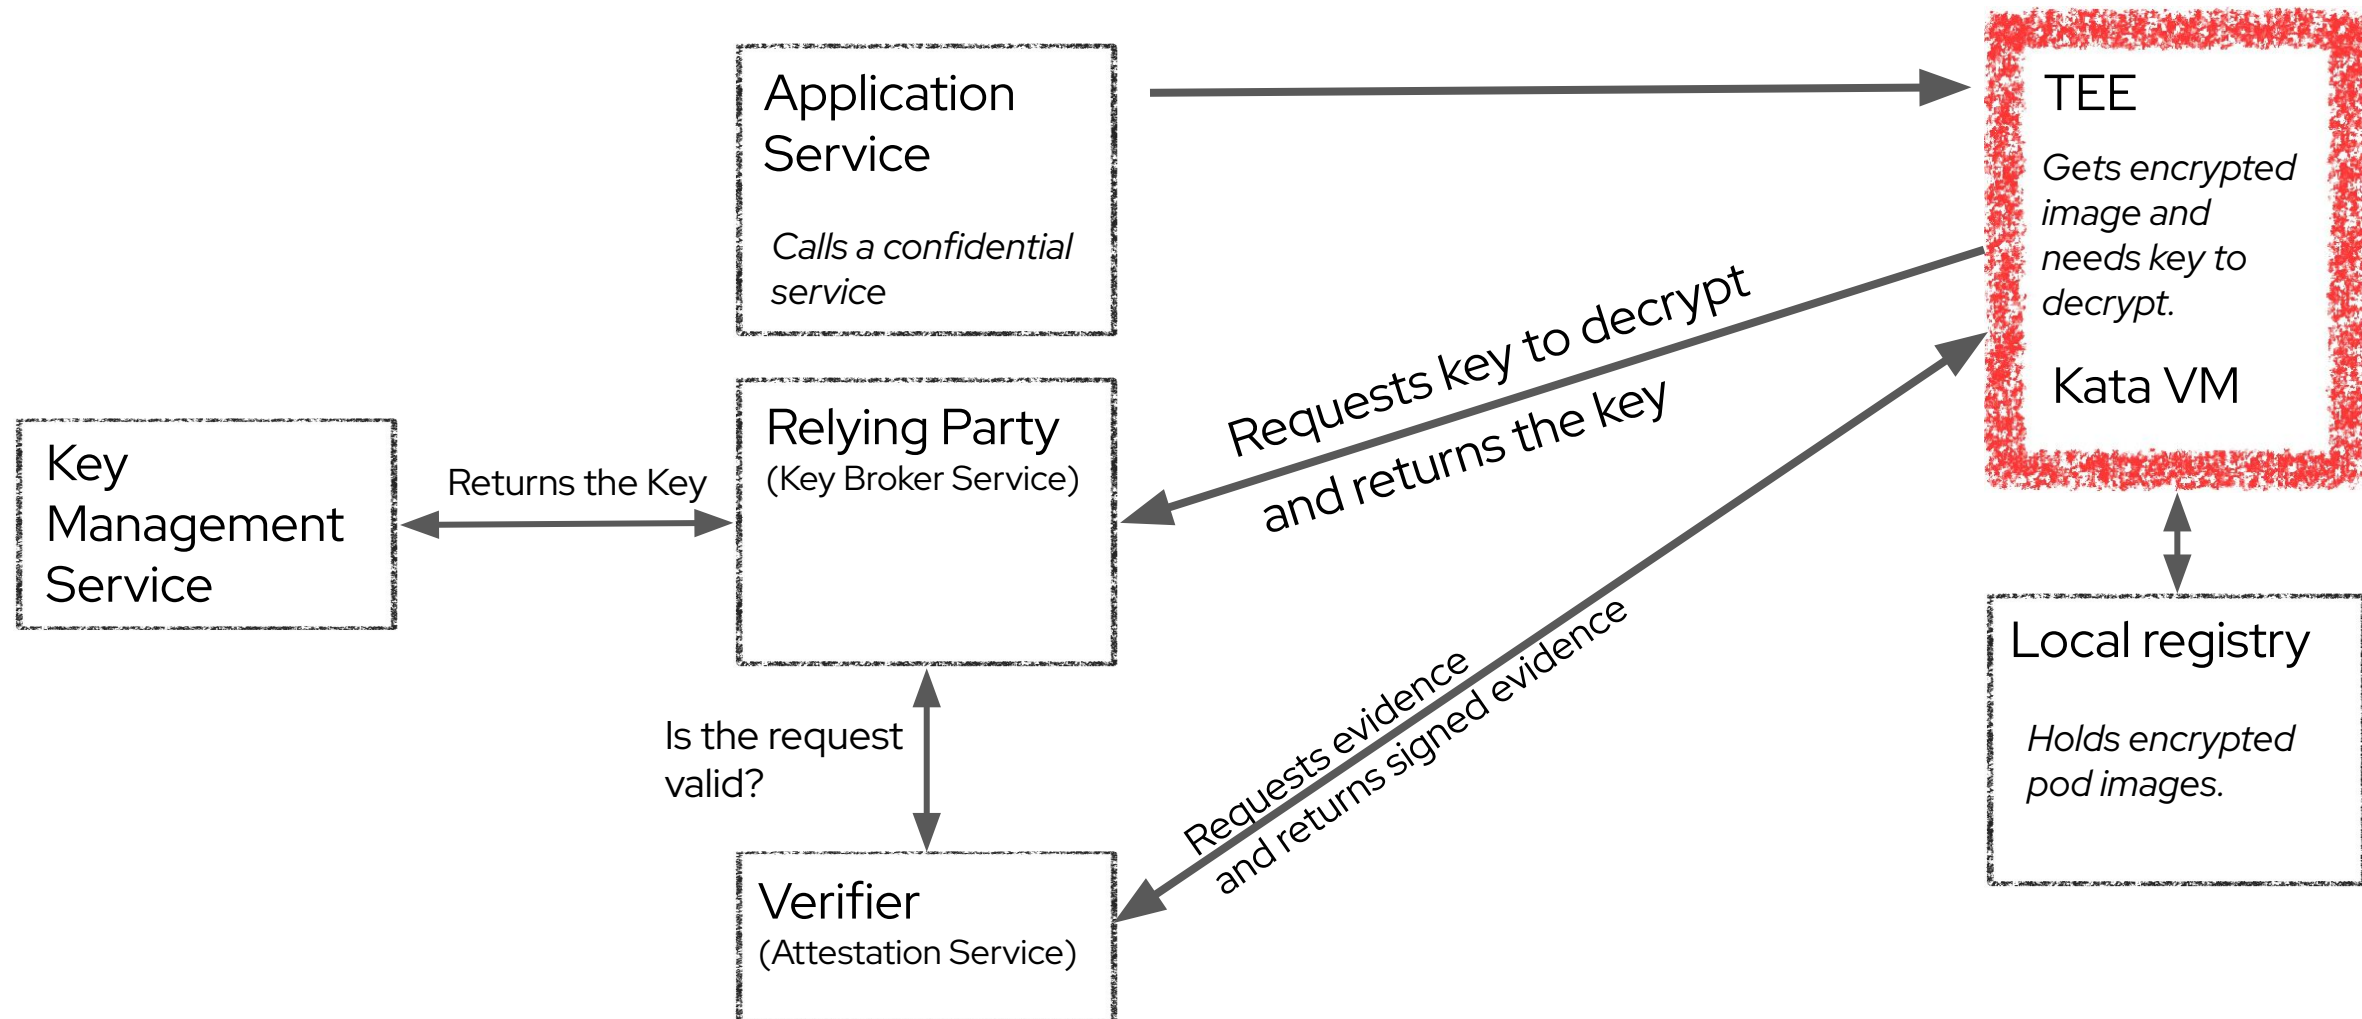
- **Key Management Service (KMS)** – A service for securely storing, managing and backing up of cryptographic keys used by applications and users.

- **Image Build Service** – Services used to build confidential containers or VM images for end users.
- **Image Registry** – A service that is used to store encrypted and/or signed container and VM images required for CC workloads. Examples of such registries include [Quay.io](#), [Docker Hub](#), CSPs provided registries, etc.

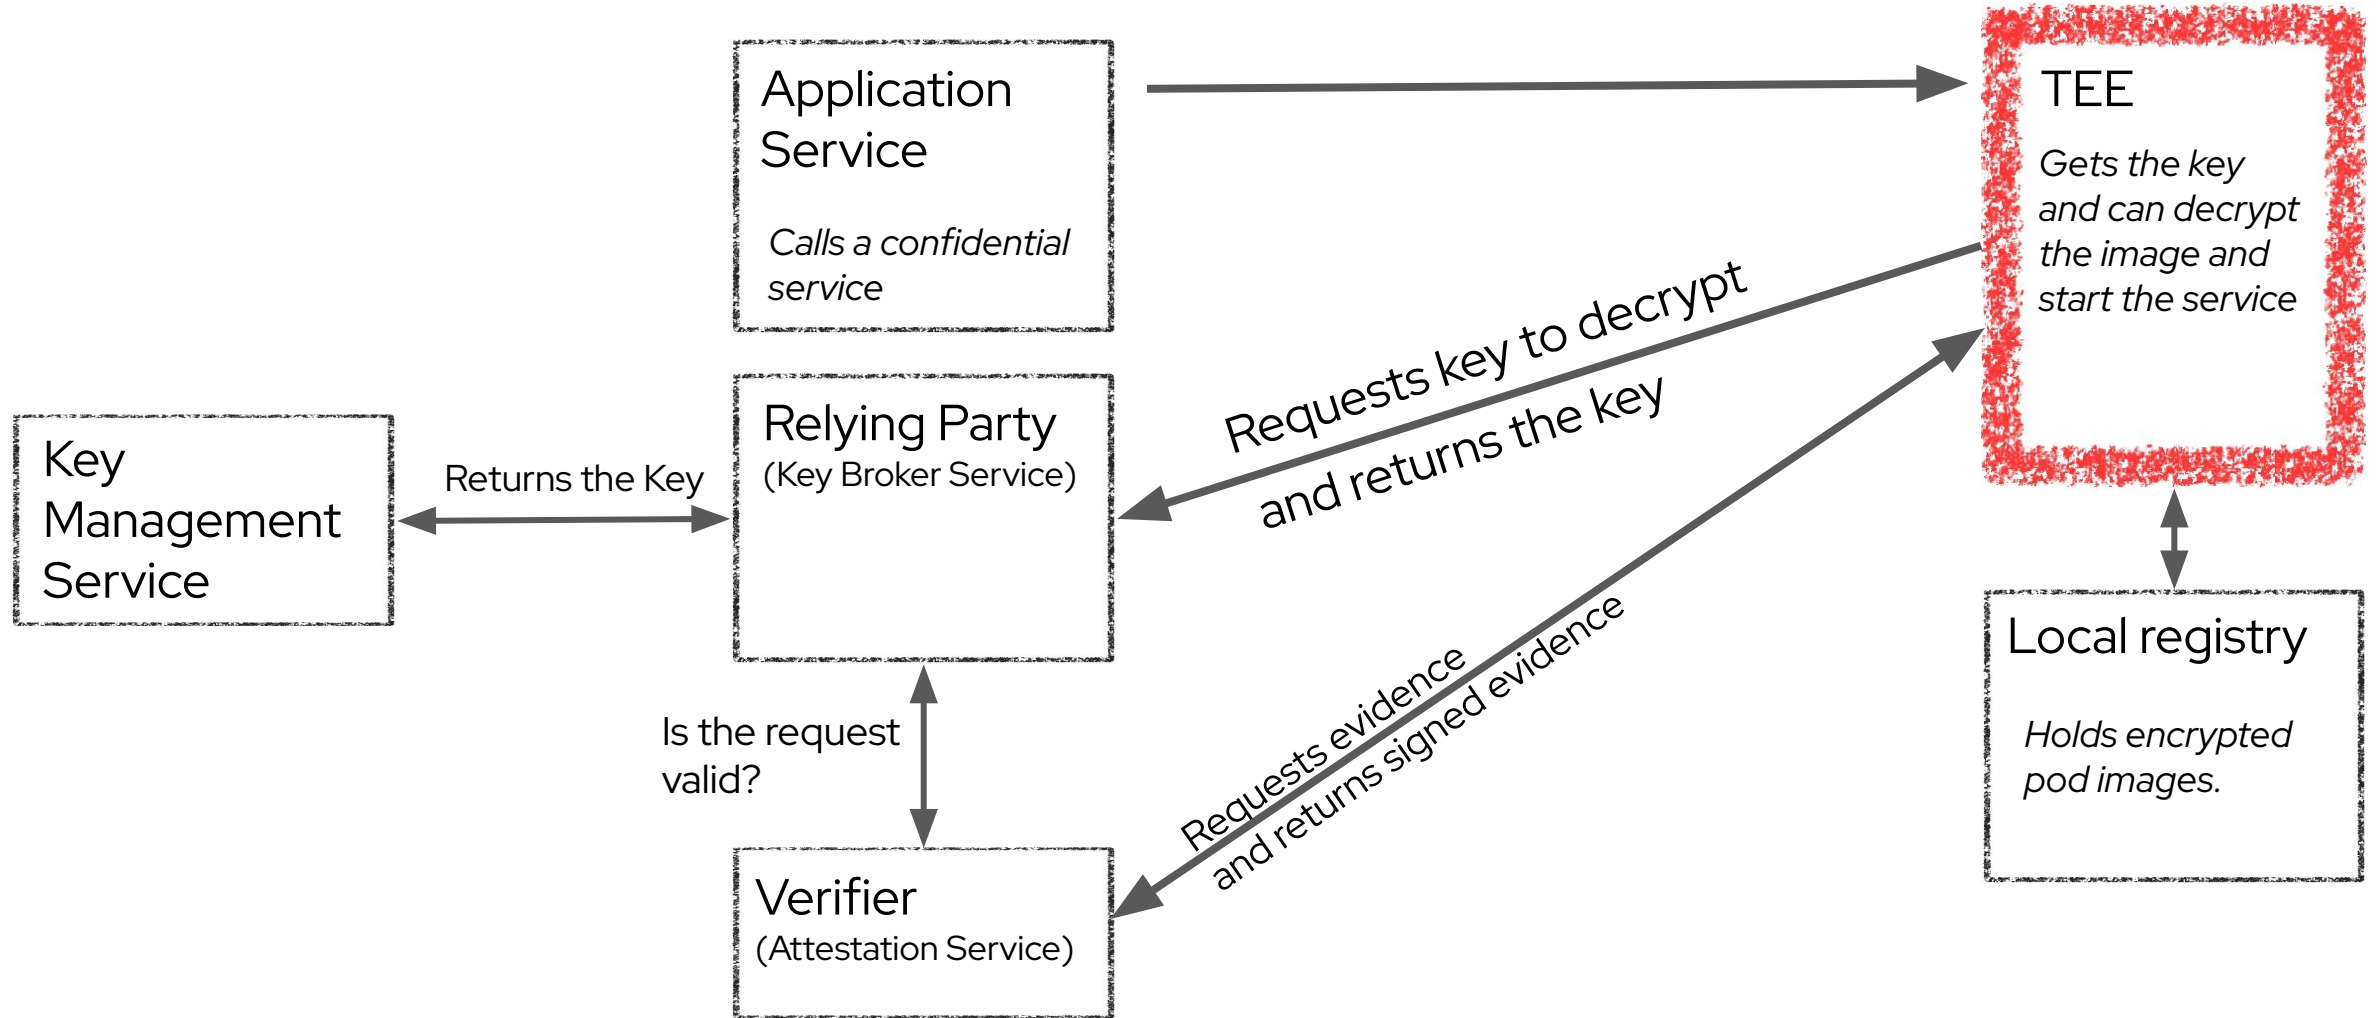
Flow of the attestation



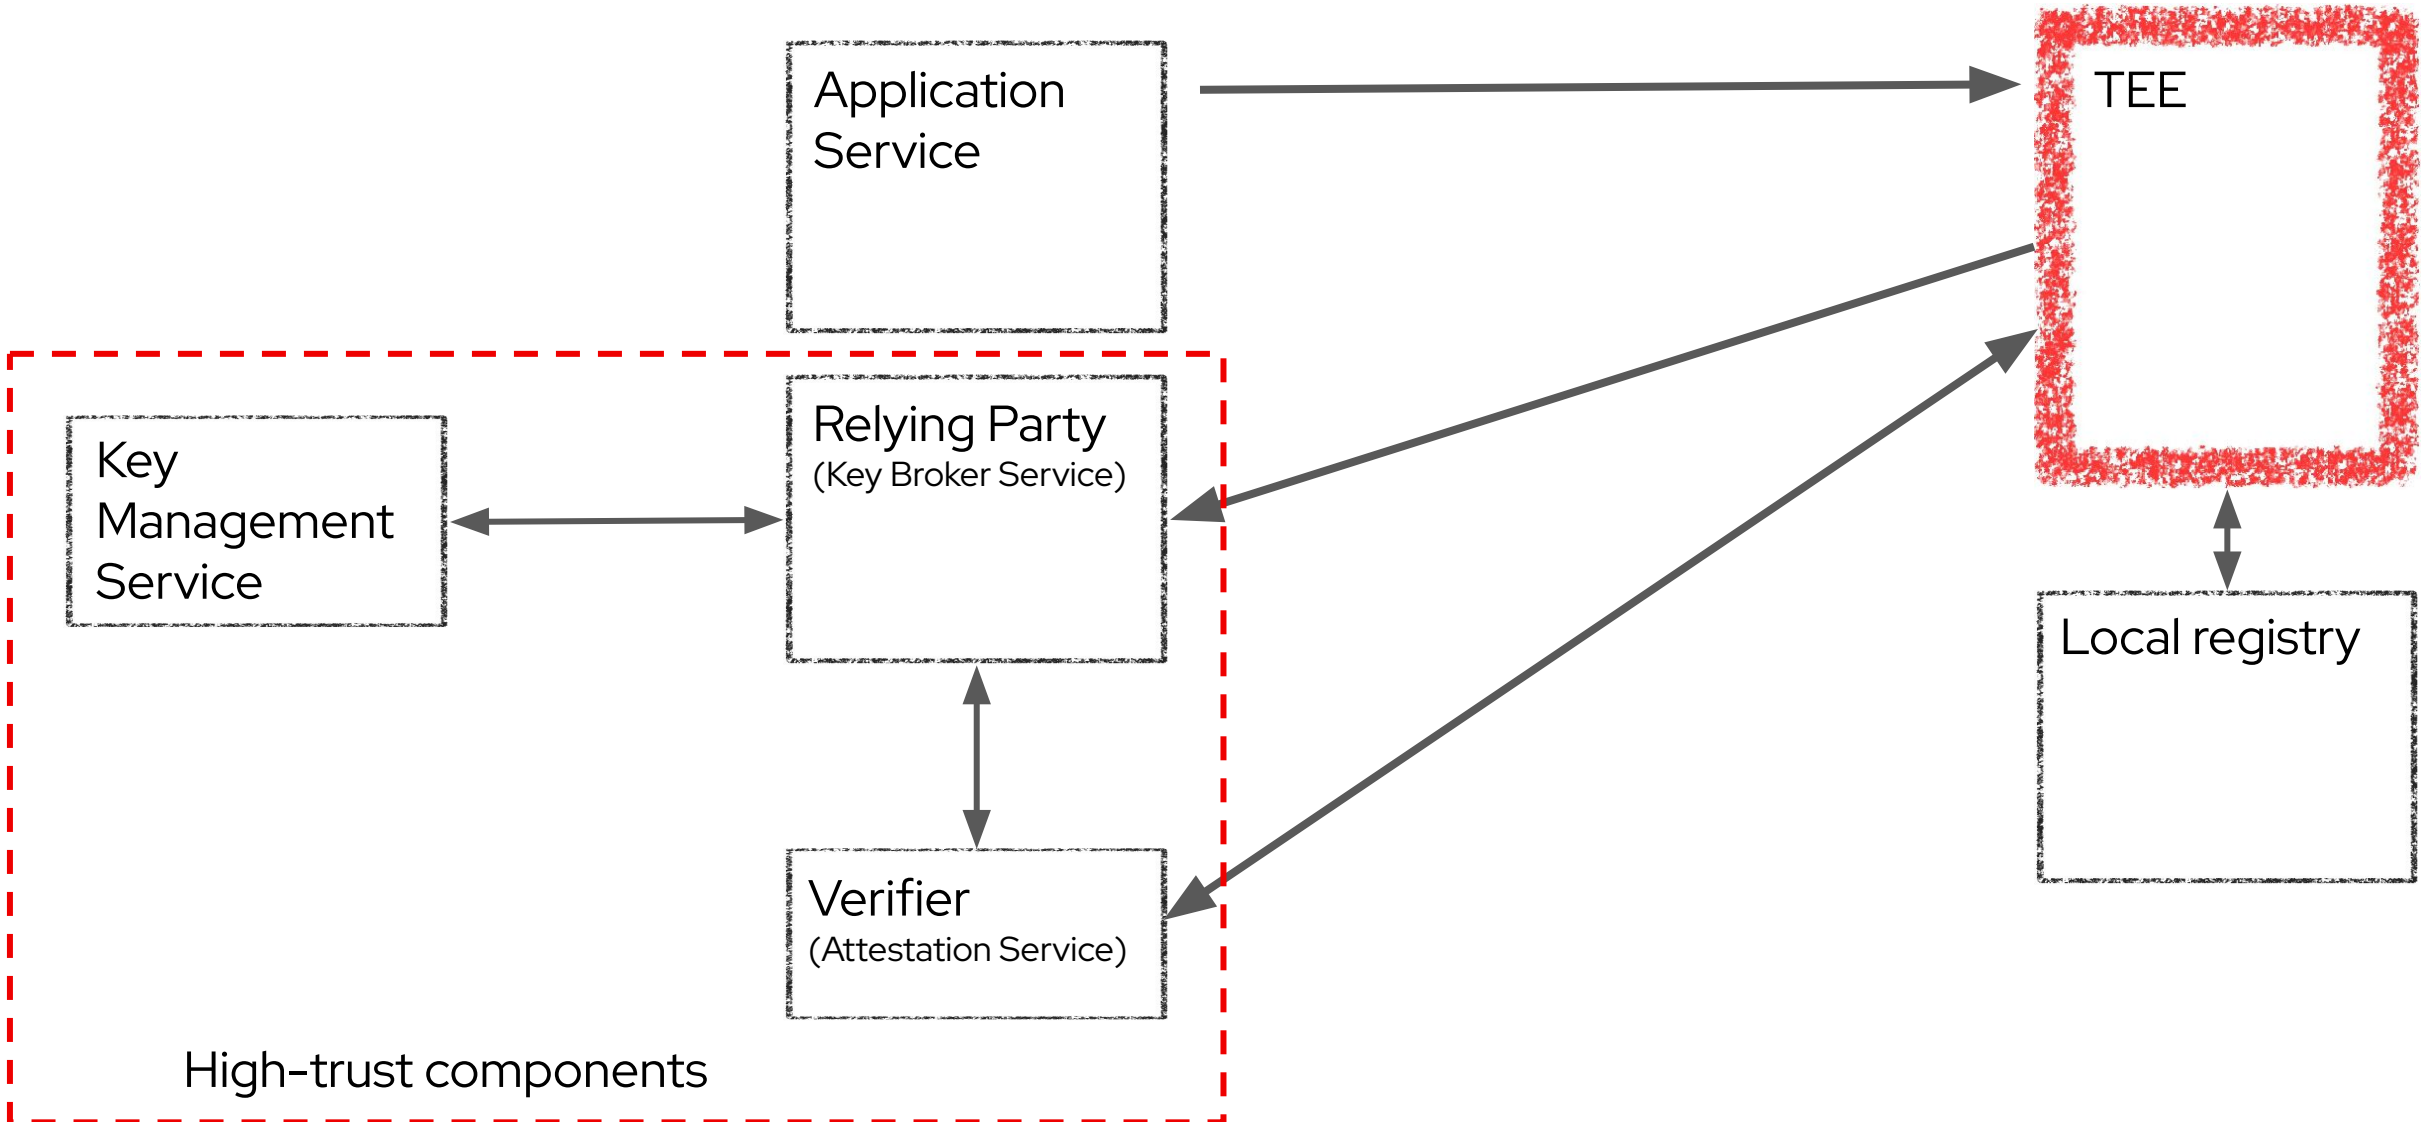
Flow of the attestation



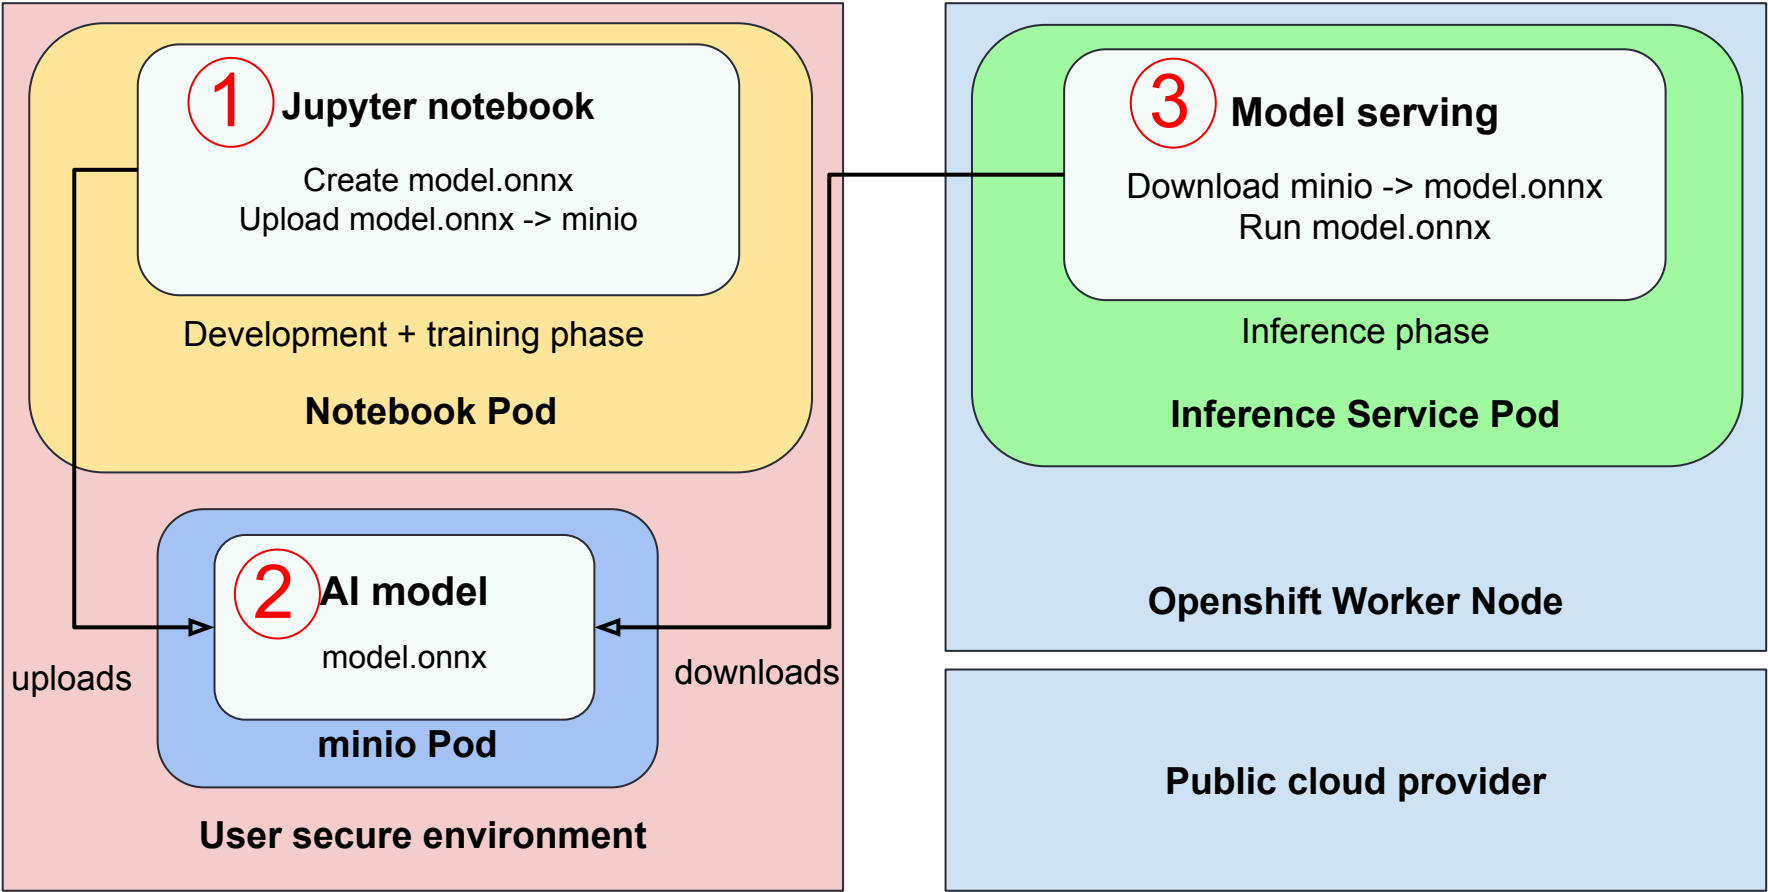
Flow of the attestation



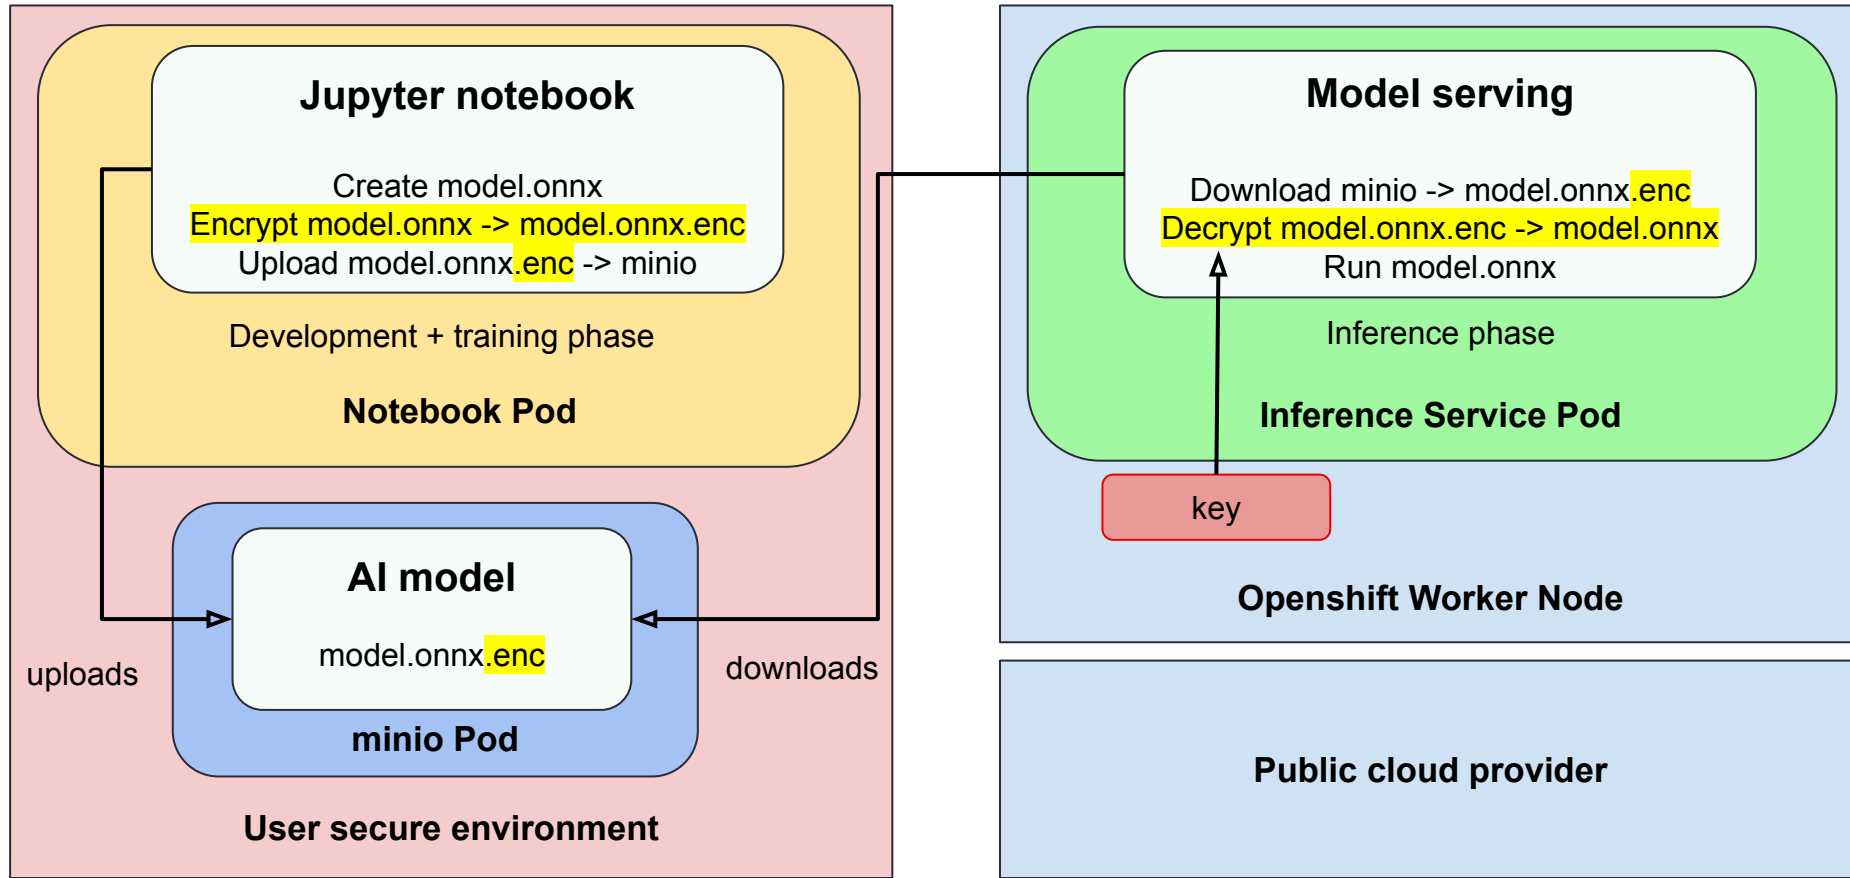
Flow of the attestation



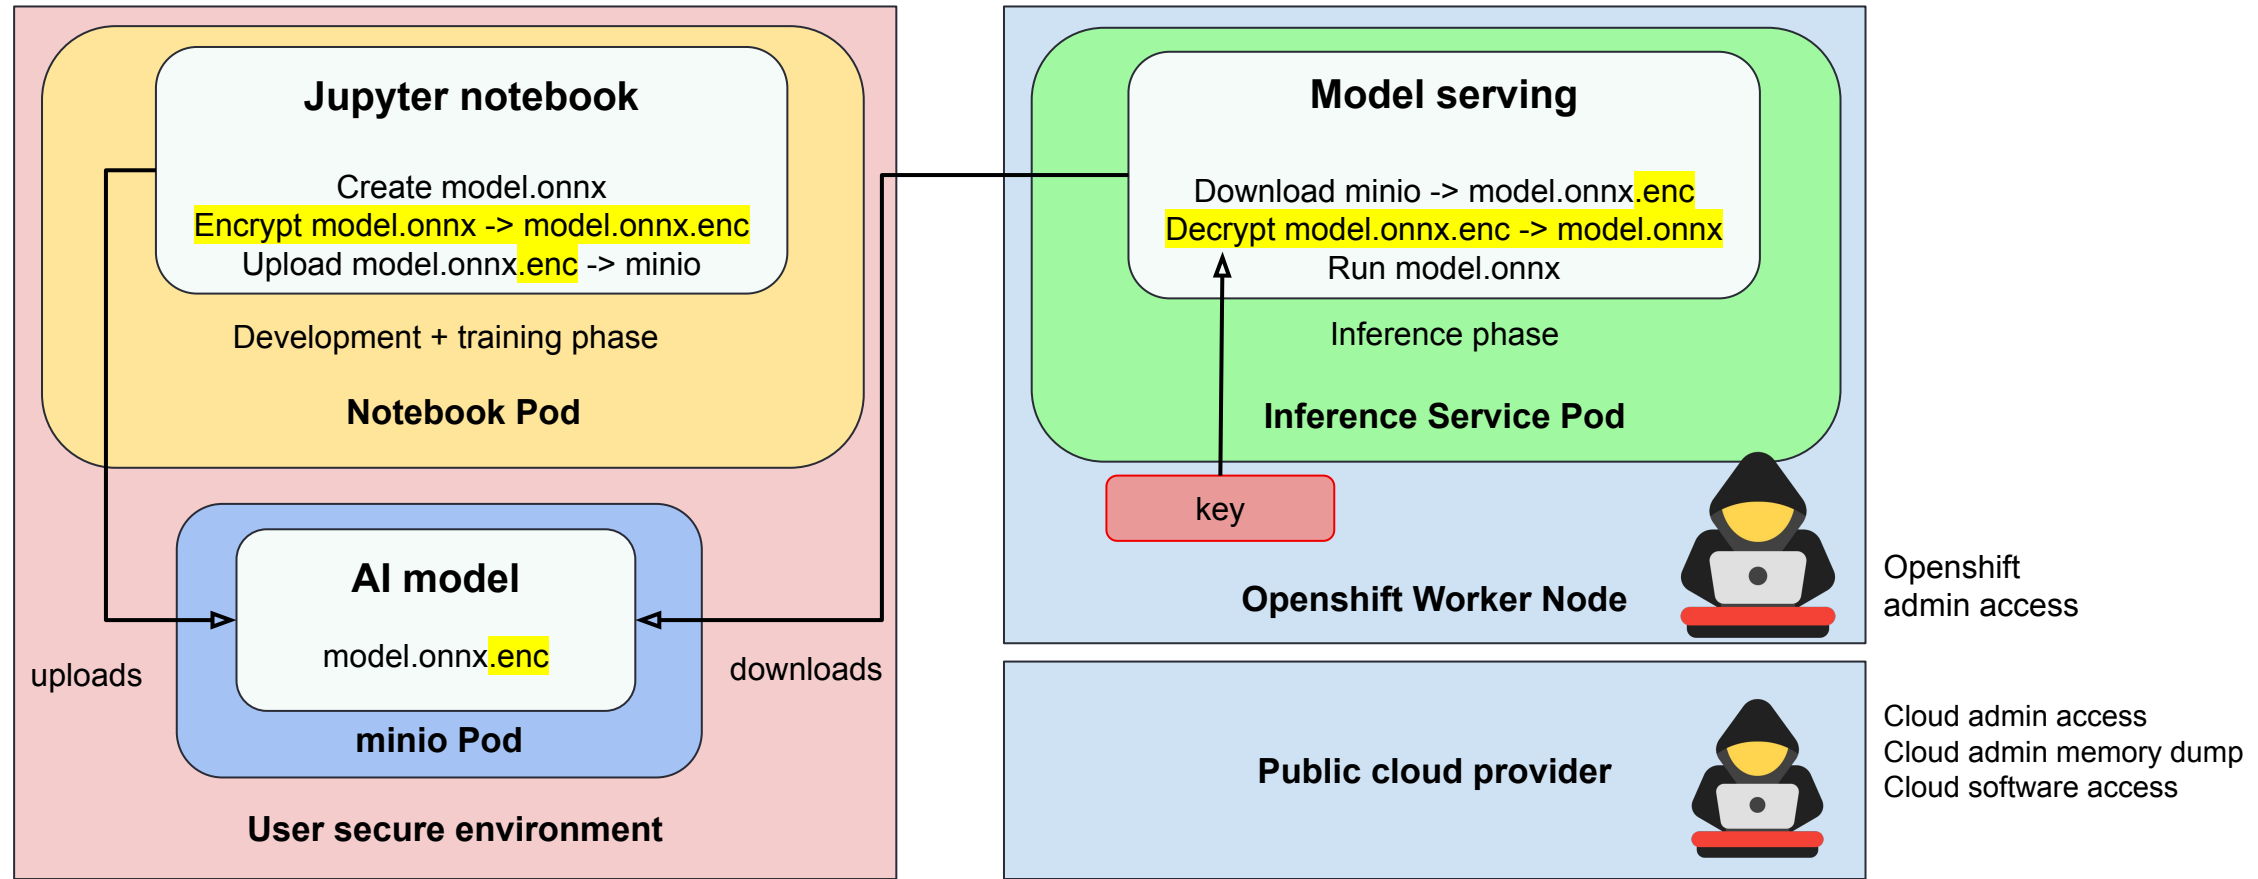
Assumptions



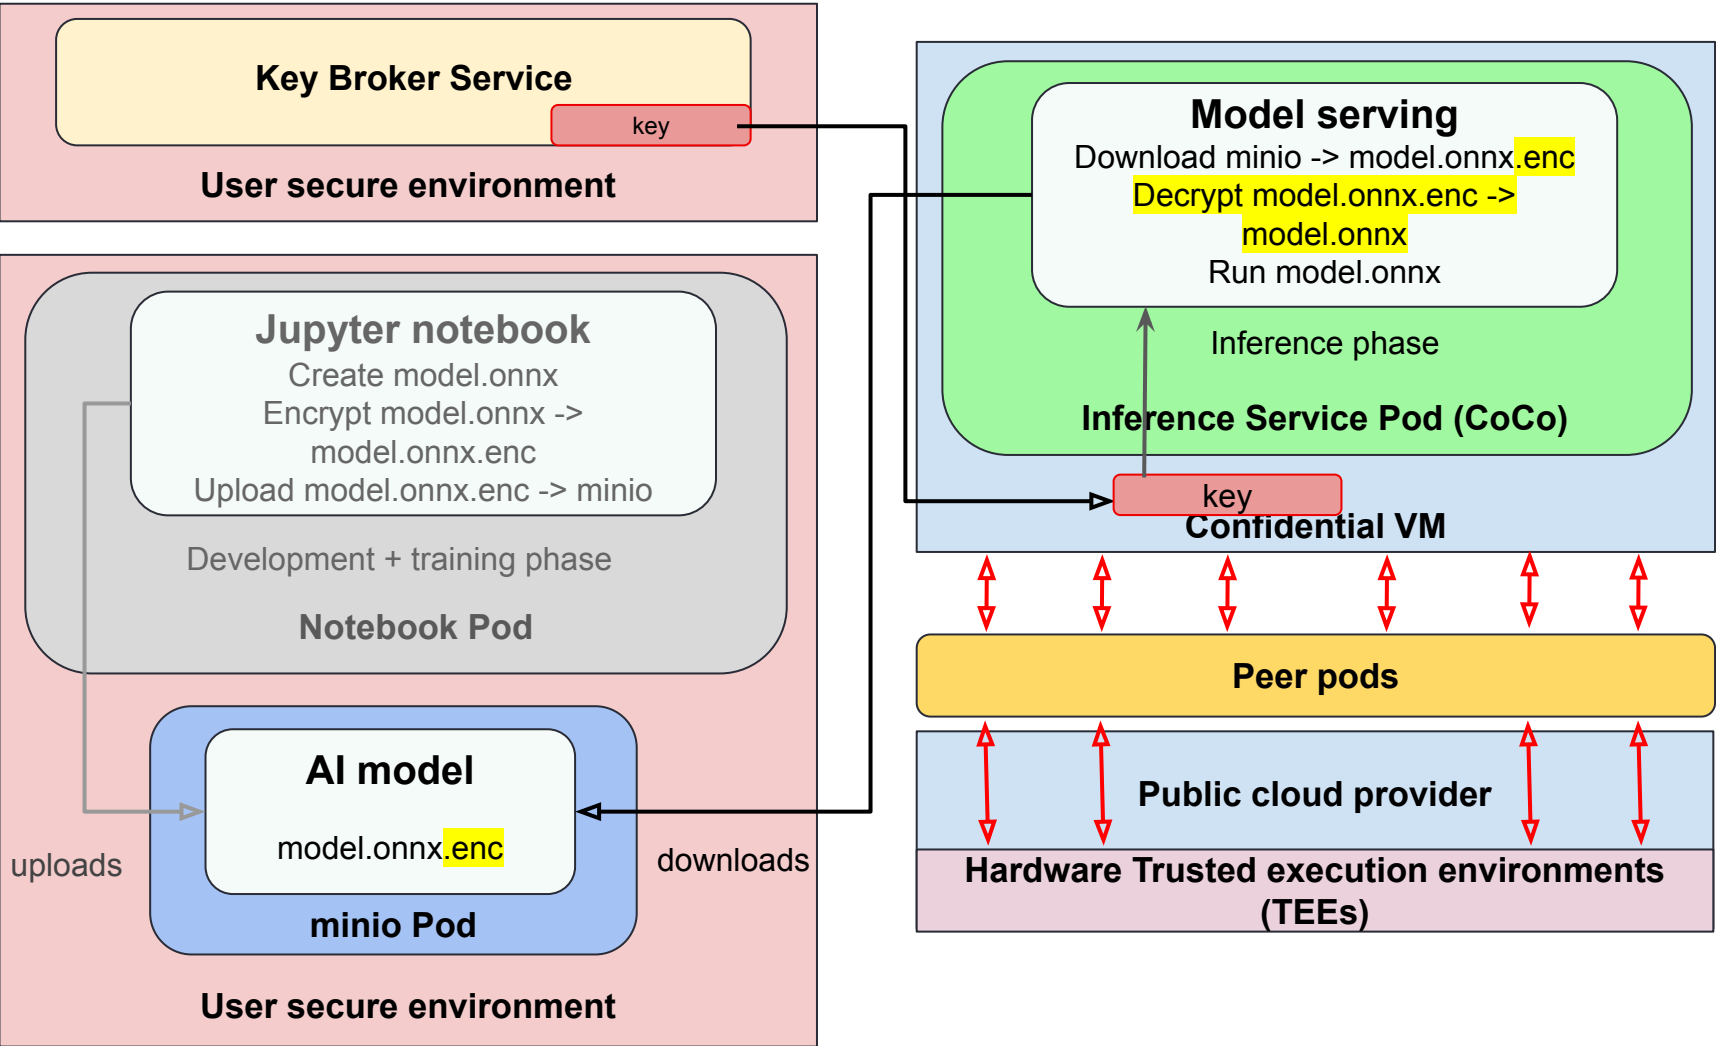
Fraud detection demo scenario with model encryption



Fraud detection demo scenario with model encryption



Fraud detection demo scenario with CoCo



Categorization of Use cases


- ▶ What use cases are our customers telling us about?

Partner Interaction




2 protected datasets interacting in confidential container

Secure Cloudburst




Using the public cloud to for peak workload or shared resources

IP Protection/Integrity




Protection of supplier data and business logic in customer environments

Edge use case




Protecting Edge device data in the public cloud for aggregation

Total Tenant Isolation




Isolating OpenShift Tenants

Digital Sovereignty

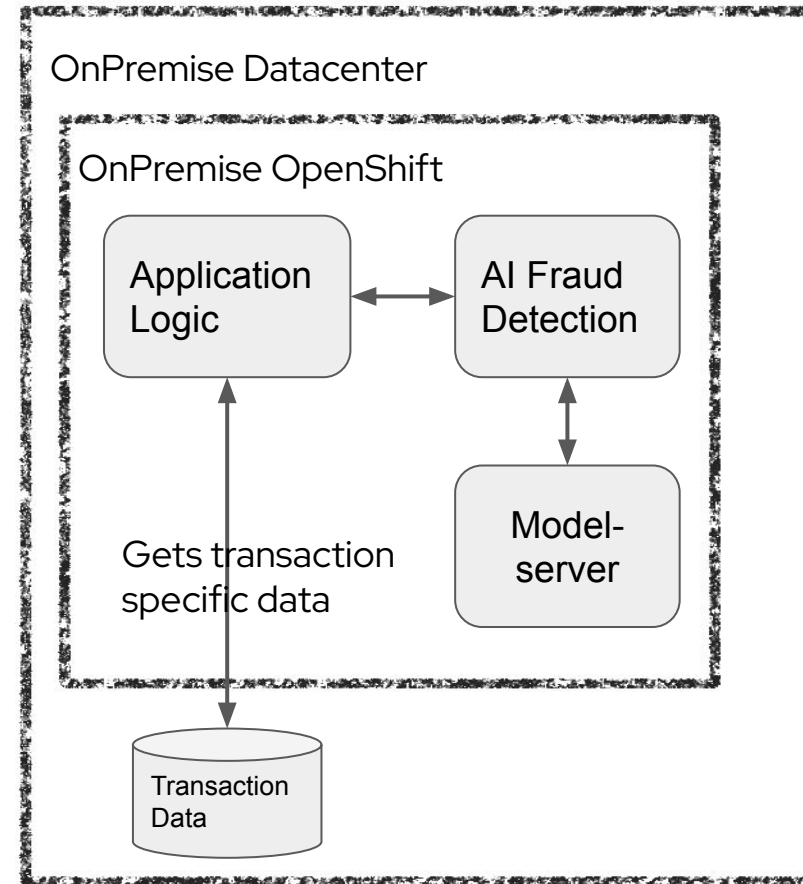


Encapsulating and moving workload from one provider to the next.

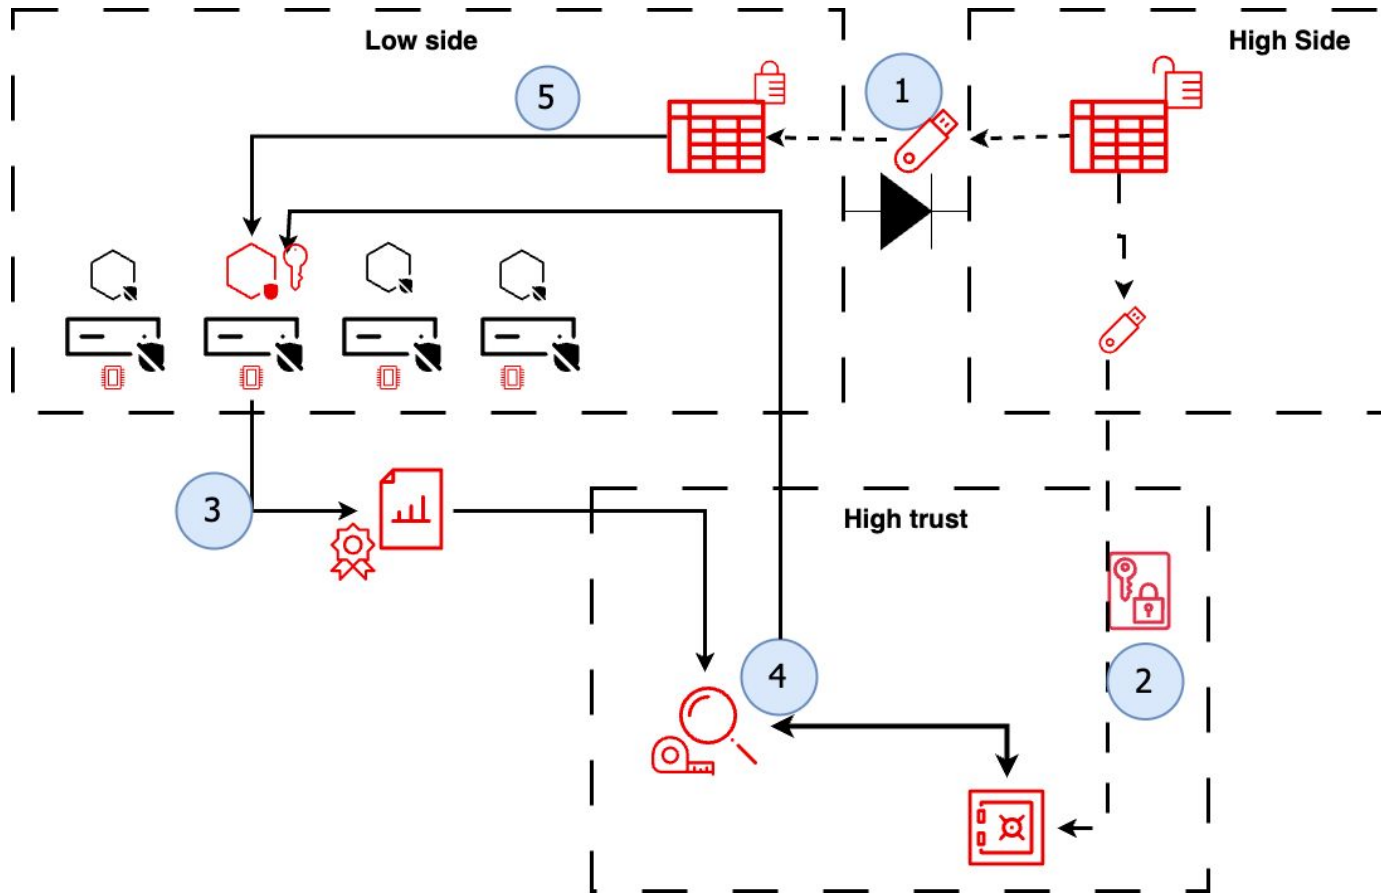


Confidential Computing Usecase (Fraud Analysis)

- Analyse transactions for fraud.
- AI model has been trained on premise.
- Transaction data and customer data is on premise.
- Goal is to use that AI model. Unfortunately company has not yet upgraded environment with GPUs.
- Usage of public cloud is difficult due to regulatory and GDPR requirements.

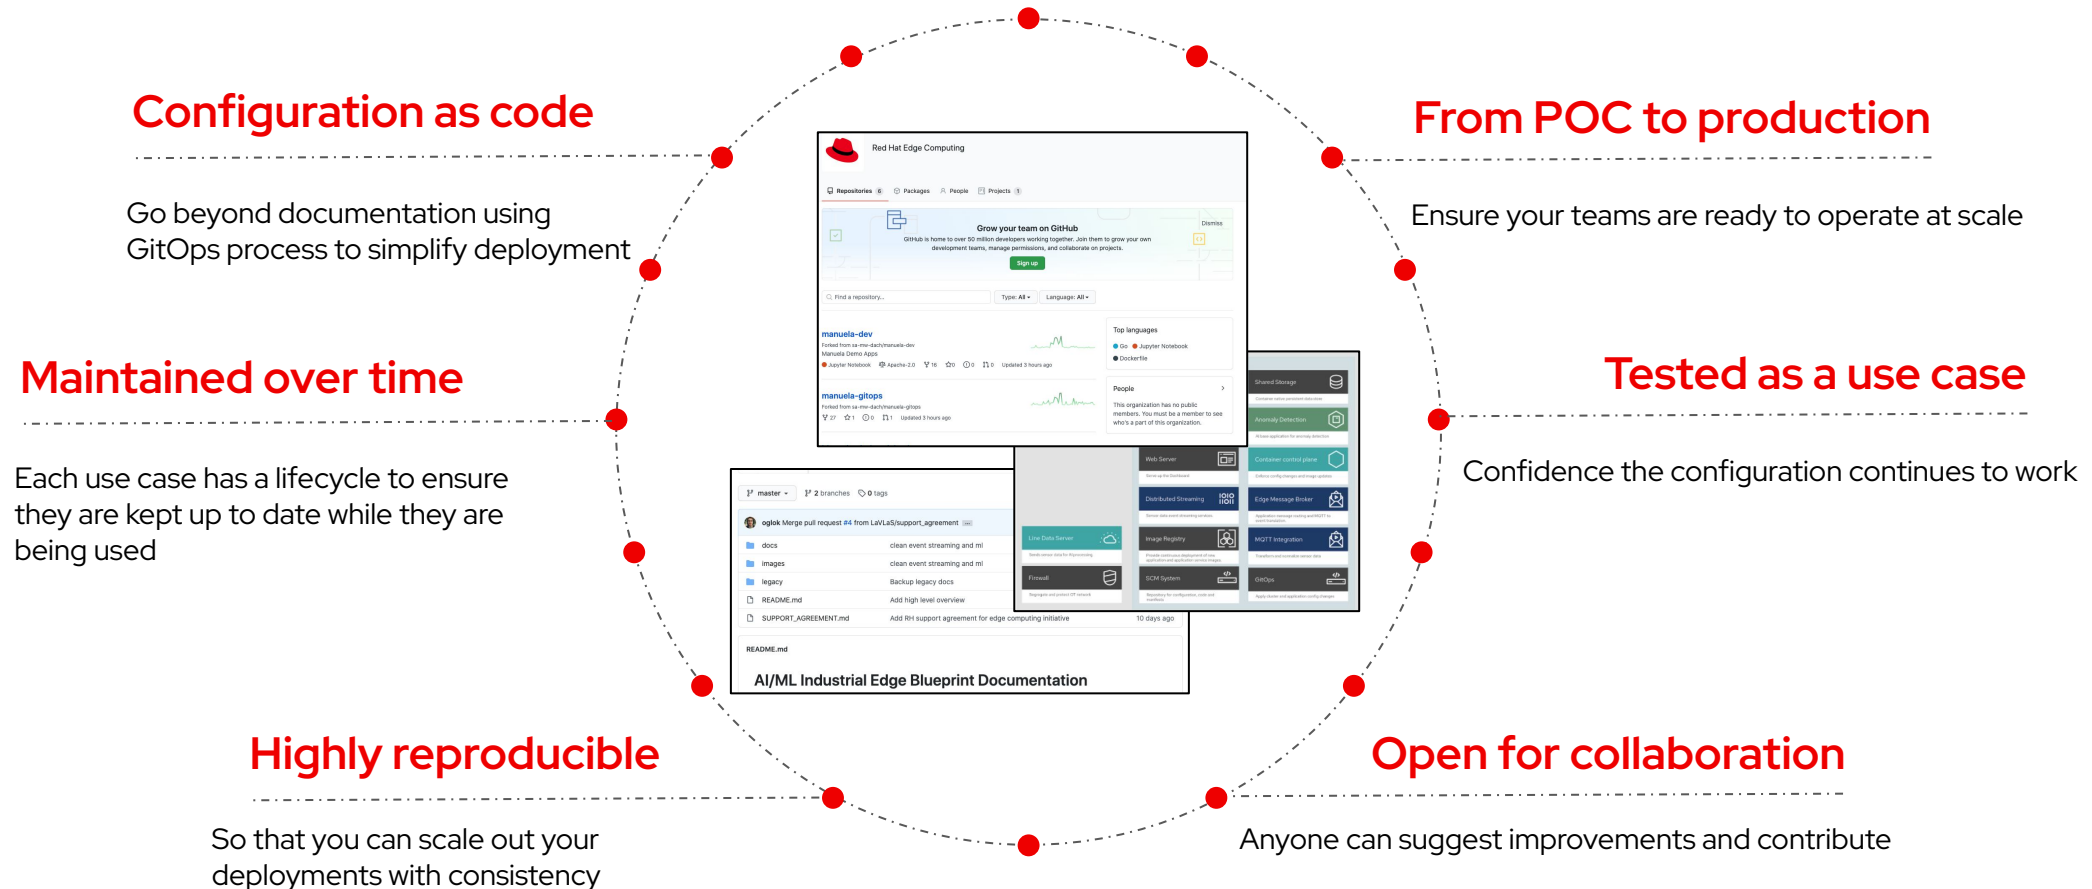


Confidential compute for 'high to low' computing



Next steps

Validated Patterns : Simplifying the creation of use cases



FAQ

- ▶ How is the kata vm integrity protected?
 - Dm-verity is used to measure rootfs integrity
- ▶ What about oc exec / oc copy etc
 - oc exec is disabled by tenant side configuration (in the enclave) which is protected by the dm-verity device map

▶

Key links

- ▶ [OpenShift Sandbox Containers \[GA\]](#)
- ▶ [Confidential Containers \[Upstream\]](#)
- ▶ [Scripts used today](#)
- ▶ [Containers used today](#)
- ▶ CoCo Blogs
 - [What is the Confidential Containers project?](#)
 - [Understanding the Confidential Containers Attestation Flow](#)
 - [CoCo architecture](#)
 - [CoCo quick start guide](#)
 - [CoCo release notes](#)

The End

Hardware implementations

- ▶ So this confidentiality needs to be implemented at the hardware level.
 - AMD SEV SNP: Confidentiality on a Core based VM and on the Memory
<https://www.amd.com/en/processors/epyc-confidential-computing-cloud>
 - Intel SGX (core based) and TDX (VM)
<https://www.intel.com/content/www/us/en/products/docs/processors/xeon-accelerated/security-accelerators-product-brief.html>
 - IBM z HyperProtect + Secure Execution
HSM modules available to enable confidential computing environments
 - ARM CCA
<https://www.arm.com/architecture/security-features/arm-confidential-compute-architecture>
 - AWS Nitro
<https://docs.aws.amazon.com/whitepapers/latest/security-design-of-aws-nitro-system>
 - RISC-V (plan)
<https://github.com/riscv-non-isa/riscv-ap-tee/blob/main/specification/riscv-cove.pdf>

Confidential Computing Organizations

► There are 2 standardizing organizations:

- CCC (founded 2019)
 - The Confidential Computing Consortium (CCC) brings together hardware vendors, cloud providers, and software developers to accelerate the adoption of Trusted Execution Environment (TEE) technologies and standards.
 - <https://confidentialcomputing.io/>

You're in good company.

accenture

蚂蚁集团
ANT GROUP

arm

Google

HUAWEI

intel

Meta

Microsoft

Red Hat

- CNCF (accepted 03/2022)

- The Cloud Native Computing Foundation (CNCf) hosts critical components of the global technology infrastructure.
<https://www.cncf.io/>



Red Hat (member)

Red Hat CNCf Platinum Member

CNCf Members · Platinum

Red Hat is a software company that offers enterprise open source software solutions.